

# Reinforcement learning

Episode 9  $\frac{3}{4}$ , 2018

## More policy gradients



Yandex  
Data Factory


LAMBDA 



**British Hedgehog  
Preservation Society**

# Continuous action spaces

- Regular MDP
- $a \in \mathbb{R}^n$




Which methods can we use?

# Continuous action spaces

- We can learn critic easily
- The problem is finding

$$a_{opt}(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$



Worst case: optimize over neural net!

# Normalized advantage functions

Idea 1: restrict  $Q(s,a)$  so that optimization becomes trivial

For example, parabola (for 1d action space)

$$Q(s, a) = V(s) + A(s, a)$$

$$A(s, a) = -k_{\theta}(s) \cdot (a - \mu_{\theta}(s))^2$$

How to find optimal  $a$ ?

# Normalized advantage functions

Idea 1: restrict  $Q(s,a)$  so that optimization becomes trivial

For example, parabola (for 1d action space)

$$Q(s, a) = V(s) + A(s, a)$$

$$A(s, a) = -k_{\theta}(s) \cdot (a - \mu_{\theta}(s))^2$$

How to find optimal  $a$ ? -  $a_{\text{opt}} = \mu(s)$

# Normalized advantage functions

Idea 1: restrict  $Q(s,a)$  so that optimization becomes trivial

For example, parabola (for 1d action space)

$$Q(s, a) = V(s) + A(s, a)$$

$$A(s, a) = -k_{\theta}(s) \cdot (a - \mu_{\theta}(s))^2$$

**Q:** How does it generalize for n-dimensional  $\mathbf{a}$ ?

# Normalized advantage functions

Idea 1: restrict  $Q(s,a)$  so that optimization becomes trivial

For example, parabola (for 1d action space)

$$Q(s, a) = V(s) + A(s, a)$$

$$A(s, a) = -0.5 \cdot (a - \mu_\theta(s))^T \cdot L(s) \cdot L(s)^T (a - \mu_\theta(s))$$

Where  $L(s)$  is a lower-triangular matrix

# Normalized advantage functions


Network:

(trains end-to-end)

$$Q(s, a) = V(s) + A(s, a)$$

$$A(s, a) = \dots$$

$$\underset{\theta}{\operatorname{argmin}} \left( Q(s_t, a_t) - [r + \gamma \cdot V(s_{t+1})] \right)^2$$





# Deterministic policy gradient

- Idea2: learn a separate network to find  $a_{opt}$
- Train critic  $Q_{\theta}(s, a)$

$$\underset{\theta}{argmin} \left( Q(s_t, a_t) - [r + \gamma \cdot V(s_{t+1})] \right)^2$$

- Train actor  $a_{opt}(s) \approx \mu_{\theta}(s)$

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

# Deterministic policy gradient

- Idea2: learn a separate network to find  $a_{opt}$

- Train critic  $Q_{\theta}(s, a)$

$$\underset{\theta}{\operatorname{argmin}} \left( Q(s_t, a_t) - [r + \gamma \cdot V(s_{t+1})] \right)^2$$

How do we get  $V(s')$ ?



- Train actor  $a_{opt}(s) \approx \mu_{\theta}(s)$

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

# Deterministic policy gradient

- Idea2: learn a separate network to find  $a_{opt}$
- Train critic  $Q_{\theta}(s, a)$


$$\underset{\theta}{argmin} \left( Q(s_t, a_t) - [r + \gamma \cdot Q(s_{t+1}, \mu_{\theta}(s_{t+1}))] \right)^2$$

- Train actor  $a_{opt}(s) \approx \mu_{\theta}(s)$

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

# Deterministic policy gradient

■ actual returns




# Deterministic policy gradient

- Gradient approximation: 
$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

■ actual returns


■ critic (approx)




# Deterministic policy gradient

- Gradient approximation:

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$



# Going neural



# Duct tape zone

- In general
  - “Natural” for continuous action spaces
    - Discrete: use gumbel-softmax, [bit.ly/2v0Xfpz](https://bit.ly/2v0Xfpz)
  - Approximation is best around current policy
  - Weak critic can introduce bias
- vs. REINFORCE
  - Better off-policy
  - Less variance if reward is smooth
  - (subjectively) harder to tune



# Deterministic policy gradient

Demo with torcs <http://bit.ly/2pXwdKa>

